



What the AI era means for cybersecurity, and how federal agencies can get ready

Frontier models are compressing cyberattack timelines and reshaping how leaders should measure and remediate risk. Federal agencies that act now will hold the defensive advantage before these capabilities reach adversaries, over roughly the next 60 days.

How the Mythos moment is reshaping federal cyber

For years, sophisticated attacks required sophisticated attackers, and defenders usually had a buffer to remediate before exploitation. That buffer is collapsing. In April 2026, Anthropic's Claude Mythos marked a step change in AI's ability to find and exploit vulnerabilities. Through Project Glasswing, select organizations can now apply Mythos-class models, including Claude Mythos 5, Fable 5, and Opus 4.8, defensively. These models don't just find flaws; they chain low-severity findings into viable exploit paths that would otherwise sit invisible in a backlog. As open-weight models reach the same frontier, that asymmetric window closes.

Federal agencies now face converging technical and regulatory pressure. OMB M-26-05 replaced compliance attestation with tailored, risk-based assurance matched to mission need; a June 2026 Executive Order directs CISA to mandate AI-enabled defensive tools through Binding Operational Directives by July 2. The capability is real, the window is finite, and the gap between policy and execution is where mission risk now concentrates.

Harden your critical assets with FedCyber.AI Preview, powered by Anthropic

A federal-grade agentic cyber harness that lets agencies apply Mythos-class models to strengthen their security posture at scale and speed, without building the capability from scratch.

THREE ACTIVE-DEFENSE OUTCOMES



Zero-day discovery

Apply frontier models to source code and applications to surface zero-day vulnerabilities before public disclosure or exploitation.




Exploit-chain identification

Map how individual lower-severity findings combine across your infrastructure into a viable exploit path no single finding reveals on its own.




Prioritized remediation routing

Correlate findings with asset criticality and proven exploitability, focusing patching capacity on what creates real mission risk, not CVSS base scores alone.



Each finding pairs with a validated fix, ready to open as a pull request.



FedCyber.AI ranks exploitable findings across your entire software estate.

WHAT YOU GET

Everything the Cyber Readiness Sprint delivers:

- ✓ **FedCyber.AI.** An agentic cyber platform that applies Mythos-class models inside your existing cloud boundary, at scale and speed.
- ✓ **Proprietary expertise.** The joint Accenture and Anthropic team already delivering agentic cyber impact for leading organizations.
- ✓ **Forward Deployed Engineers.** A scalable team that validates AI findings and accelerates remediation of confirmed high-priority vulnerabilities.
- ✓ **Actionable intelligence.** Identified exploit chains, a prioritized remediation backlog, and a tailored software-assurance baseline.

THE CYBER READINESS SPRINT · 4-8 WEEKS

AFS and Anthropic deploy FedCyber.AI to identify exploitable chains, close the most critical exposure gaps, and stand up the tailored software-assurance baseline that M-26-05 and the CISA directives require. Forward Deployed Security Engineers validate AI findings and embed alongside your engineers to build lasting internal capability.

GET STARTED

Start your Cyber Readiness Sprint

Visit AFS.com



Matthew Finney

Managing Director,
Anthropic Lead



Amanda Satterwhite

Managing Director,
Cyber Lead